

Endianness and ARM® Processors

Overview

The terms little-endian and big-endian refer to the way in which a word of data is stored into sequential bytes of memory. The first byte of a sequence may store either the least significant byte of a word (little-endian) or the most significant byte of a word (big-endian). Typically, this is a detail that is of no relevance to the software engineer. In certain cases, however, the software engineer must be careful to take the endianness of the hardware into account.

ARM® processors use 32-bit words¹, and if the software accesses all data as 32-bit words, the issue of endianness is not relevant. However, if the software executes instructions that operate on data 8 or 16 bits at a time, and that data must be mapped at specific memory addresses (such as with memory-mapped I/O), then the issue of endianness will arise. The endianness of the system as a whole is implementation dependent and is determined by the circuitry that connects the processor to its peripheral components.

Programs that will execute directly from flash memory must be stored into flash in the correct endian format so that the code will be fetched correctly by the processor. For this reason, the Target Flash Programming feature of SourcePoint provides the user with the option to reverse the endianness of the flash code before programming it into a flash device.

What Is Endianness?

Endianness refers to how bytes and 16-bit half-words map to 32-bit words in the system memory. Given that a 32-bit word contains 4 bytes or 2 half-words, two possibilities exist for the ordering of bytes and half-words within the word. Suppose that a program must deal with a hexadecimal number of 0xBBCCDDEE. Figure 1 illustrates how that number appears in a register:

Bits 31-20	Bits 19-16	Bits 15-08	Bits 07-00
BB	CC	DD	EE

Figure 1: Register Layout

If the memory location where this value is stored is displayed in 32-bit word format, the number will appear exactly as it appears in the register. When that number is accessed as bytes or half-words, the order of the sub-fields depends on the endian configuration of the system. If a program stores the above value at location 0x100 as a word and then fetches the data as individual bytes, two possible orders exist.

In the case of a little-endian system, the data bytes will have the order depicted by **Figure 2**. Note that the rightmost byte of the word is the first byte in the memory location at 0x100. This is why this format is called little-endian; the least significant byte of the word occupies the lowest byte address within the word in memory.

¹ When not operating in Thumb™ mode

Endianness and ARM® Processors

Address	0x100	0x101	0x102	0x103
Data	0xEE	0xDD	0xCC	0xBB

Figure 2: Little Endian Byte Order

If the program executes in a big-endian system, the word in **Figure 1** has the following byte order in memory:

Address	0x100	0x101	0x102	0x103
Data	0xBB	0xCC	0xDD	0xEE

Figure 3: Big Endian Byte Order

The least significant byte of the word is stored in the high order byte address. The most significant byte of the word occupies the low order byte address, which is why this format is called big-endian. Operations with 16-bit half-words have similar consequences.

When dealing with half-words, the memory address must be a multiple of two. Thus the value in Figure 1 will occupy two half-word addresses: 0x100 and 0x102. Figure 4 shows the layout for both endian configurations.

Address	0x100	0x102
Little Endian Data	0xDDEE	0xBBCC
Big Endian Data	0xBBCC	0xDDEE

Figure 4 Half Word Endian Orders

Note: Within the half-word, the bytes maintain the same order as they have in the word format. In little-endian mode, the least significant half-word resides at the low-order address (0x100) and the most significant half-word resides at the high-order address (0x102). For the big-endian case, the layout is reversed.

Usually the issue of endianness is transparent to both programmers and users. However, the issue becomes relevant when data must cross between endian formats, such as may be the case for networks or distributed systems that use different platforms. The next section discusses how endianness relates to ARM processors.

Ramifications of Endianness and ARM Processors

The [ARM Architectural Reference Manual](#) states that ARM processors are bi-endian, meaning they can operate in either little-endian or big-endian modes. The ARM processor does not have any instructions or features that affect endianness. The endianness of the system as a whole is determined by the circuitry that connects the processor to its peripheral devices.

In those ARM cores that contain a System Control coprocessor (coprocessor 15), Control Register 1 contains a bit that can be used to reverse the endian configuration of the system. The overall design of the system may require that this bit be set for proper operation. Check the user's manual for the board to determine the proper setting for this bit.

Since this bit is cleared at reset, the software engineer still must ensure that any code that executes either directly from flash or immediately after reset is formatted for the correct endian configuration, and the reset handler code must set this bit if required.

For those ARM cores that do not include coprocessor 15, the software engineer has no control over the endian

Endianness and ARM® Processors

configuration of the system.

ARM specifications state that data values accessed in word format are invariant with respect to endianness. If a program stores a 32-bit value at a given memory address, then switches the endian configuration of the processor and reads back the 32-bit value at that same address, it will get the correct result. However, if data are read or written in smaller chunks (8 or 16 bits), this will no longer hold true. See "An Endian Experiment" below for some examples of what happens if the endian configuration is changed for 8- or 16-bit data.

Endian Issues When Programming Flash Devices

Today's flash memory devices are typically 8 or 16 bits wide. Some systems may implement a 32-bit-wide flash memory interface, but typically this actually consists of two interleaved 16-bit devices. Programming operations on these devices involve 8- or 16-bit data write operations at specific addresses within each device. For this reason, the software engineer must know and understand the endian configuration of the hardware in order to successfully program the flash device(s). This information typically is documented in the hardware reference manual for the board.

There are two main factors that must be considered to correctly program a flash device:

A flash programming operation is initiated by placing the flash device into a special mode. This typically is accomplished by writing an 8- or 16-bit value to a particular address within the device. Where this address is mapped in the processor memory space will depend on the endian configuration of the system.

Code which will be executed directly from an 8- or 16-bit flash device must be stored in a way that instructions will be properly recognized when they are fetched by the processor. This may be affected by the endian configuration of the system. Compilers typically have a switch that can be used to control the endianness of the code image that will be programmed into the flash device.

If the Target Flash Programming feature of SourcePoint is used to program the flash device, the user may select the

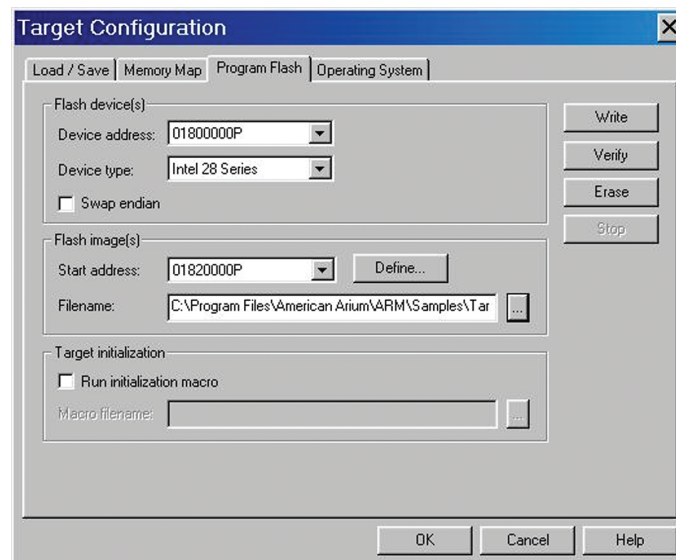


Figure 5: SourcePoint Program Flash tab on the Target Configuration dialog box

Endianness and ARM® Processors

“Swap Endian” option to reverse the endianness of the code image.

When selecting a device type from the "Device type" pull-down menu (Figure 5), the user must specify the endian configuration of the target. Currently, this menu provides the following choices:

- Intel 28 Series
- Intel 28 Series – Big Endian
- AMD 29 Series
- AMD 29 Series – Big Endian
- SST 39 Series
- SST 39 Series – Big Endian

If the endian configuration of the target does not match what the user has selected (e.g., big endian or not) the flash programming dialog will report that it was unable to locate or identify the flash device.

Also note that in many cases the user must provide a pre-flash initialization script to initialize the hardware into a state where a flash programming operation can be performed successfully.

An Endian Experiment

This section shows the results generated by a SourcePoint command file that writes an incrementing data pattern to target memory starting at address zero. The data are written to memory as 32-bit words, with the target in little endian mode. The data are then displayed in word, half word, and byte formats. The system is then switched to big endian mode and the data are displayed again. The entire experiment is then repeated with the target initially in big endian mode and switching to little endian mode.

Note that the byte order within each word does not change when the endian configuration is changed after data are written. However, the order of bytes within each word is affected by what endian mode is in effect when the data are written.

```
>//Load data in little endian mode
>control.b = 0
>load C:\ArmProjects\RomFiles\IncByteFilled.bin at 0
>//Display data in word format
>ord4 0 len 10
00000000 03020100 07060504 0B0A0908 0F0E0D0C
00000010 13121110 17161514 1B1A1918 1F1E1D1C
00000020 23222120 27262524 2B2A2928 2F2E2D2C
00000030 33323130 37363534 3B3A3938 3F3E3D3C
>//Display data in half word format
>ord2 0 len 20
00000000 0100 0302 0504 0706 0908 0B0A 0D0C 0F0E
00000010 1110 1312 1514 1716 1918 1B1A 1D1C 1F1E
00000020 2120 2322 2524 2726 2928 2B2A 2D2C 2F2E
00000030 3130 3332 3534 3736 3938 3B3A 3D3C 3F3E
>ord1 0 len 40
>//Display data in byte format
00000000 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F
00000010 10 11 12 13 14 15 16 17 18 19 1A 1B 1C 1D 1E 1F
00000020 20 21 22 23 24 25 26 27 28 29 2A 2B 2C 2D 2E 2F
00000030 30 31 32 33 34 35 36 37 38 39 3A 3B 3C 3D 3E 3F
>//Switch to big endian mode
>control.b=1
```

Endianness and ARM® Processors

```

>//Display data in word format
>ord4 0 len 10
00000000 03020100 07060504 0B0A0908 0F0E0D0C
00000010 13121110 17161514 1B1A1918 1F1E1D1C
00000020 23222120 27262524 2B2A2928 2F2E2D2C
00000030 33323130 37363534 3B3A3938 3F3E3D3C
>//Display data in half word format
>ord2 0 len 20
00000000 0302 0100 0706 0504 0B0A 0908 0F0E 0D0C
00000010 1312 1110 1716 1514 1B1A 1918 1F1E 1D1C
00000020 2322 2120 2726 2524 2B2A 2928 2F2E 2D2C
00000030 3332 3130 3736 3534 3B3A 3938 3F3E 3D3C
>//Display data in byte format
>ord1 0 len 40
00000000 03 02 01 00 07 06 05 04 0B 0A 09 08 0F 0E 0D 0C
00000010 13 12 11 10 17 16 15 14 1B 1A 19 18 1F 1E 1D 1C
00000020 23 22 21 20 27 26 25 24 2B 2A 29 28 2F 2E 2D 2C
00000030 33 32 31 30 37 36 35 34 3B 3A 39 38 3F 3E 3D 3C
>//Load data in big endian mode
>load C:\ArmProjects\RomFiles\IncByteFilled.bin at 0
>//Display data in word format
>ord4 0 len 10
00000000 00010203 04050607 08090A0B 0C0D0E0F
00000010 10111213 14151617 18191A1B 1C1D1E1F
00000020 20212223 24252627 28292A2B 2C2D2E2F
00000030 30313233 34353637 38393A3B 3C3D3E3F
>//Display data in half word format
>ord2 0 len 20
00000000 0001 0203 0405 0607 0809 0A0B 0C0D 0E0F
00000010 1011 1213 1415 1617 1819 1A1B 1C1D 1E1F
00000020 2021 2223 2425 2627 2829 2A2B 2C2D 2E2F
00000030 3031 3233 3435 3637 3839 3A3B 3C3D 3E3F
>//Display data in byte format
>ord1 0 len 40
00000000 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F
00000010 10 11 12 13 14 15 16 17 18 19 1A 1B 1C 1D 1E 1F
00000020 20 21 22 23 24 25 26 27 28 29 2A 2B 2C 2D 2E 2F
00000030 30 31 32 33 34 35 36 37 38 39 3A 3B 3C 3D 3E 3F
>//Switch to little endian mode
>control.b = 0
>//Display data in word format
>ord4 0 len 10
00000000 00010203 04050607 08090A0B 0C0D0E0F
00000010 10111213 14151617 18191A1B 1C1D1E1F
00000020 20212223 24252627 28292A2B 2C2D2E2F
00000030 30313233 34353637 38393A3B 3C3D3E3F
>//Display data in half word format
>ord2 0 len 20
00000000 0203 0001 0607 0405 0A0B 0809 0E0F 0C0D
00000010 1213 1011 1617 1415 1A1B 1819 1E1F 1C1D
00000020 2223 2021 2627 2425 2A2B 2829 2E2F 2C2D
00000030 3233 3031 3637 3435 3A3B 3839 3E3F 3C3D
>//Display data in byte format
>ord1 0 len 40
00000000 03 02 01 00 07 06 05 04 0B 0A 09 08 0F 0E 0D 0C
00000010 13 12 11 10 17 16 15 14 1B 1A 19 18 1F 1E 1D 1C
00000020 23 22 21 20 27 26 25 24 2B 2A 29 28 2F 2E 2D 2C

```

Endianness and ARM® Processors

00000030 33 32 31 30 37 36 35 34 3B 3A 39 38 3F 3E 3D 3C

Summary

In many cases the endian configuration of a system is not relevant to the software engineer. However, in systems that contain memory-mapped peripheral devices (such as flash memory devices), endianness must be considered whenever the software must access data in bytes or half-words. See the user's manual for your board to determine the correct endian configuration.

The Target Flash Programming feature of SourcePoint allows the user to select the proper endian configuration of the target system before initiating a flash programming operation.

An initialization macro may be required to initialize the board into a state where it is ready for flash programming. This initialization macro may also need to set the endian control bit in the System Control Register (see the user's manual for your board).

