

# **64-bit Linux Device Driver Porting**

John A. Ronciak  
Staff Engineer  
Intel Corp.

January 21 - 25, 2002

# Agenda

- **IA-32 Linux Device Drivers**
- **Make the Device Driver 64-bit Code Clean**
- **Building a 64-bit Linux Device Driver**
- **Debugging a 64-bit Linux Device Driver**
- **Code Examples ( Case Studies )**
- **Additional Information**
- **Summary**

## IA-32 Linux Device Drivers

# Current Device Drivers

- Use a current IA-32 known working driver
- Fully test the IA-32 driver and device on known working IA-32 hardware
  - Known working tools
  - Everybody understands IA-32 opcodes (assembly code)
  - More machines around for development use (more resources)

## IA-32 Linux Device Drivers

# Current Device Drivers

- **System Hardware Requirements:**
  - **IA-32 Machine**
    - **550 MHz Intel® Pentium® III Processor**
    - **128MB RAM**
    - **1 4-GB Hard Drive**

# Linux Environment

- **Standard Linux install**
  - Linux Distribution of your Choice
- **Test the install for functionality**
  - Need to know things are working
- **Build your own Kernel**
  - Test again

# Standard Kernel Build

- **Get a Standard Kernel** ([www.kernel.org](http://www.kernel.org))
- **Get the Matching Kernel Debugger**  
([oss.sgi.com/projects/kdb](http://oss.sgi.com/projects/kdb))
- **Apply kdb Patch to New Kernel**
- **Build New Kernel**
- **Test New Kernel**
- **Test Debugger Functionality**
- **Test your Device Driver**

# Agenda

- IA-32 Linux Device Drivers
- **Make the Device Driver 64-bit Code Clean**
- Building a 64-bit Linux Device Driver
- Debugging a 64-bit Linux Device Driver
- Code Examples ( Case Studies )
- Additional Information
- Summary

64-bit Code Clean

# Source Code Cleaning

- **Generic 64-bit Coding Considerations**
  - Data Models
  - Data Types
  - Pointers
  - Compiler Switches
- **64-bit Driver Source Cleanup**
- **Regression Test “Cleaned” Driver on IA-32**

**Easy Porting Procedures**



64-bit Code Clean

# Generic IA-64 Considerations

- **Data Models**

- **Uniform Data Model**

- Same source runs on IA-32 and 64-bit systems

- **LP64 Data Model**

- longs and pointers are 64-bit

- ints remain 32-bit

- IA-32 Systems are ILP32

**Easy Porting Procedures**

64-bit Code Clean

# Generic 64-bit Considerations

- **Data Types**

- **New Data Types**

- `int64_t` and `uint64_t`
    - `int32_t` and `uint32_t`
    - `int16_t` and `uint16_t`
    - `int8_t` and `uint8_t`

**Easy Porting Procedures**

64-bit Code Clean

# Generic 64-bit Considerations

- **Pointers**

- All Pointers will be system defined

- i.e. 32-bit pointers for IA-32, 64-bit for IA-64

- **Do not cast pointers to int's or long's**

- Write “clean” C code

**Easy Porting Procedures**

64-bit Code Clean

# 64-bit Driver Cleaning

- Use the new data types
  - IA-32 `ulong`'s to `uint32_t`
    - `long`'s are now 64-bit
  - `ushort`'s to `uint16_t`
  - Check all hardware specific types and convert them accordingly

**Easy Porting Procedures**

# Agenda

- IA-32 Linux Device Drivers
- Make the Device Driver 64-bit Code Clean
- **Building a 64-bit Linux Device Driver**
- Debugging a 64-bit Linux Device Driver
- Code Examples ( Case Studies )
- Additional Information
- Summary

64-bit System

# 64-bit Linux Install

- **Install 64-bit Linux Distribution of your choice**
- **Configure the System as Needed**
- **menuconfig Options (xconfig)**
  - Processor Type must be Itanium (normally set)
  - Enable the Kernel debugger
    - Integrated into most Distributions
- **Build the New Kernel**

Building the 64-bit System

# 64-bit Linux System

- Reboot the New 64-bit Linux System
- Put the IA-32 Device Driver Source on the 64-bit System
- Compile both the Linux Kernel and Driver on the 64-bit System with the Native Tools
  - May need to Static Link into the Kernel
  - Pay Attention to Any and All Warnings

**Ready for Porting Now**

The 64-bit Driver

# The New Driver

- Reboot the System with the New Linux Kernel
- Test the New Driver
  - Start with basic functionality
  - Add more and more stress to the driver

**Ready for Porting Now**



The 64-bit Driver

# Uniform Data Model Driver

- **After Validation on 64-bit Platforms**
  - Regression Test the Driver on IA-32
  - Remember Uniform Data Model
  - Same Driver Must Work on Both Archs

**Easy Porting Procedures**

# Milestones so far

- **Driver Cleanly Compiles for 64-bit Linux**
- **Driver Source is validated on IA-32**
- **Driver is integrated into 64-bit Linux kernel**
- **Loadable Module support as needed**

**Ready for Porting Now**

Debugging the Driver

# Linux Debug Tools

- **Standard Printk's still work**
  - Be careful with 64-bit pointers
- **Debugger Built into the Linux Kernel**
  - It's able to:
    - Dump and modify memory
    - Disassemble code
    - Set breakpoints

**Get an Early Start, Don't Wait!**

## Debugging the Driver

# Kernel Debugger

- Some Debugger commands:
  - ***md*** and ***mds*** displays memory
    - [kdb\kdb\\_md.html](kdb\kdb_md.html)
  - ***mm*** modifies memory
    - [kdb\kdb\\_mm.html](kdb\kdb_mm.html)
  - ***id*** disassembles code
    - [kdb\kdb\\_bp.html](kdb\kdb_bp.html)
  - ***rd*** and ***rm*** displays and modifies registers
    - [kdb\kdb\\_rd.html](kdb\kdb_rd.html)

**Get an Early Start, Don't Wait!**

## Debugging the Driver

# Kernel Debugger (cont.)

- **bt** and **btp** stack backtrace and for pid
  - [kdb\kdb\\_bt.html](#)
- **bp(bph)**, **bl**, **bpa**, **bc**, **be**, **bd** set and manipulate breakpoints
  - [kdb\kdb\\_bp.html](#)
- **ss** and **ssb** single step and step to branch
  - [kdb\kdb\\_ss.html](#)
- **env** and **set** show and set env variables
  - [kdb\kdb\\_env.html](#)

**Get an Early Start, Don't Wait!**

Debugging the Driver

# Kernel Debugger (cont.)

- **cpu** switch to new cpu
- **reboot** will immediately reboot machine
- **go** continues execution
- **ef** display exception frame

**Get an Early Start, Don't Wait!**

Debugging the Driver

# Kernel Debugger (cont.)

- Use the Debugger
  - Saves development time
  - Use to test error/unused code paths
  - Enabled for panic backtraces
  - Use to check for correct types

**Get an Early Start, Don't Wait!**

# Agenda

- **IA-32 Linux Device Drivers**
- **Make the Device Driver 64-bit Code Clean**
- **Building a 64-bit Linux Device Driver**
- **Debugging a 64-bit Linux Device Driver**
- **Code Examples ( Case Studies )**
- **Additional Information**
- **Summary**



## Case Studies

# Code Examples

- **Structure packing**
  - Compiler generates codes with naturally aligned boundaries.
- **Structures use hard-coded size for padding**

```
Struct Buffer {  
    void *Ptr[10];  
    char  Padding[88];  
}
```

To pad data to 128-byte chunk (  $4 \times 10 + 88$  ) in IA-32

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

- Structures Padding ( cont.)
  - Pointer is 8-bytes in 64-bit code! Better to use  
`Char Padding[128 - (10 * sizeof(void *))]`
- Performance impact
  - due to padding error

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

- Don't use Unaligned Data Accesses in the kernel
- IT IS NOT ALLOWED!!!
  - It will cause a panic!
  - Check to make sure your data is aligned!
  - It is allowed in user space but will impact performance and should still be avoided

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

```
Struct {  
    ulong Space;  
    void *Buffer;  
    uint Offset;  
    ulong Length;  
} IoBlock;
```

- **Is there any thing wrong in the example above? Will this work in 64-bit Linux?**

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

```
typedef struct _cb_header_t {  
    ushort    cb_status;        /* Command Block Status*/  
    ushort    cb_cmd;           /* Command Block Command*/  
    ulong     cb_lnk_ptr;       /* Link To Next CB*/  
} cb_header_t, *pcb_header_t __attribute__((__packed__));
```

New IA-64 structure:

```
typedef struct _cb_header_t {  
    uint16_t   cb_status;        /* Command Block Status*/  
    uint16_t   cb_cmd;           /* Command Block Command*/  
    uint32_t   cb_lnk_ptr;       /* Link To Next CB*/  
} cb_header_t, *pcb_header_t __attribute__((__packed__));
```

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

From:

```
printk( "rfpd = 0x%x, rfd_status=0x%x\n", rfdp, rfd_status );  
printk( "    cb_cmd = 0x%x, cb_lnk_ptr = 0x%x\n",  
        rfdp->rfd_header.cb_cmd, rfdp->rfd_header.cb_lnk_ptr );
```

to:

```
printk( "rfpd = 0x%p, rfd_status=0x%x\n", rfdp, rfd_status );  
printk( "    cb_cmd = 0x%x, cb_lnk_ptr = 0x%x\n",  
        rfdp->rfd_header.cb_cmd,  
        (uint32_t)rfdp->rfd_header.cb_lnk_ptr );
```

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

From:

```
#define E1000_READ_REG(reg) \
    ((Adapter->MacType >= MAC_LIVENGOOD)? \
    readl(&((PE1000_REGISTERS)Adapter->HardwareVirtualAddress)->reg) : \
    readl(&((POLD_REGISTERS)Adapter->HardwareVirtualAddress)->reg))
```

To:

```
#define E1000_READ_REG(reg) \
    ((Adapter->MacType >= MAC_LIVENGOOD)? \
    readl((unsigned long) \
    &((PE1000_REGISTERS)Adapter->HardwareVirtualAddress)->reg) : \
    readl((unsigned long) \
    &((POLD_REGISTERS)Adapter->HardwareVirtualAddress)->reg))
```

**Get an Early Start, Don't Wait!**

## Case Studies

# Code Examples ( cont. )

- Be Aware of OS I/F Changes
- In network drivers, make sure the SKB that is handed up the stack is aligned. The IP layer will panic if the data in the SKB is not aligned using the `skb_reserve` command. This could require a copy if your hardware can't DMA to byte alignment
- Spinlocks are not changed and operate just like IA-32. Don't change them

**Get an Early Start, Don't Wait!**



# Agenda

- IA-32 Linux Device Drivers
- Make the Device Driver 64-bit Code Clean
- Building a 64-bit Linux Device Driver
- Debugging a 64-bit Linux Device Driver
- Code Examples ( Case Studies )
- Additional Information
- Summary

# Additional Information

- Useful URLs
- SDV preparation for 64-bit Linux installation
- Installing Turbo Linux

# Useful URLs and References

- **Intel® Itanium® Processor information from Intel**
  - <http://developer.intel.com/design/itanium/index.htm>
- **For kernel debugger information on (IA-32 and 64-bit Linux)**
  - <http://oss.sgi.com/projects/kdb/>
- **The Official Home Page of 64-bit Linux**
  - <http://www.linuxia64.org/>

# Useful URLs Continued

- **Where is the latest kernel source?**
  - <http://www.kernel.org/pub/linux/kernel/ia64/>
- **GNUPro tools for 64-bit Linux**
  - <http://www.redhat.com/software/tools/gnupro/>
- **First Released Distribution – TurboLinux**
  - <ftp://ftp.turbolinux.com/pub/ia64/>
  - <ftp://ftp.redhat.com/pub/redhat/redhat-7.2-en/os/ia64/>

# Boot From CD-ROM and Installing Linux

- **At EFI prompt switch to appropriate device**
  - For instance FS3: (this is your CD-ROM)
    - (assuming 2 hard disks)
- **Begin installation process**
  - elilo linux (typical)
    - LS-120 = /dev/hda
    - CD-ROM = /dev/hdc

# Installing Linux cont...

- **FDISK 3 Partitions (if not using auto)**
  - Partition 1, 300M (FAT 16)
  - Partition 2, 2000+M (Linux)
  - Partition 3, 1024M (Linux Swap)

**NOTE:** These are typical and you can make them whatever you want. The Swap space should be set to twice the amount of memory you have in the system.

# Agenda

- **IA-32 Linux Device Drivers**
- **Make the Device Driver 64-bit Code Clean**
- **Building a 64-bit Linux Device Driver**
- **Debugging a 64-bit Linux Device Driver**
- **Code Examples ( Case Studies )**
- **Additional Information**
- **Summary**

# In Summary

- **Porting Tools are Available Now**
- **The Porting Procedures are Easy**
- **Ready for the Porting Process Now**
- **Get an Early Start, Don't Wait**