

Overview of the CPU Frequency Scaling API

Introducing the CPU “Governor” Concept

Bill Gatliff

`bgat@billgatliff.com`

Freelance Embedded Systems Developer

CPU Frequency Scaling

Adjust the CPU speed:

- Increase for more processing throughput
- Decrease to lower power consumption
- Even more effective with voltage scaling

CPU Frequency Scaling

```
struct cpufreq_governor  
struct cpufreq_policy  
struct cpufreq_driver
```

CPU Frequency Scaling

“Governors” adjust speed:

- Per user request, or
- For maximum for performance, or
- For maximum battery life, or
- To match battery life to performance

But what defines “best”?

CPU Frequency Scaling

Adjust the CPU speed:

The question is, “to where?”

```
# ls /sys/devices/system/cpu/cpu0/cpufreq
# ls /sys/devices/system/cpu/cpu1/cpufreq
...
```

Governors

“Performance”

- Sets the processor to the highest allowable speed

“Powersave”

- Sets the processor to the lowest allowable speed

Governors

“Ondemand”

- Increase processor speed with scheduler loading

“Conservative”

- Same as “ondemand”, but ramps speed changes instead of stepping

Governors

“Userspace”

- Allow root users to specify processor speed

`/sys/.../cpufreq/scaling_setspeed`

- Userspace governor only
- Shows current CPU speed
- Write to change CPU speed
- Subject to rounding and policy limits

“What Governors Do I Have?”

`/sys/.../cpufreq/scaling_available_governors`

- Lists available governors

`/sys/.../cpufreq/scaling_governor`

- Name of currently active governor
- Write a new name to change
- (Not all governors are supported on all CPUs)

CPU Frequency Policy

`/sys/.../cpufreq/scaling_min_freq`

`/sys/.../cpufreq/scaling_max_freq`

- Maximum and minimum limits (policy)

`/sys/.../cpufreq/cpuinfo_min_freq`

`/sys/.../cpufreq/cpuinfo_max_freq`

- Maximum and minimum limits (hardware)

CPU Frequency Policy

`/sys/.../cpufreq/scaling_driver`

- Name of the driver implementing scaling

struct cpufreq_driver

Adjusts CPU frequency:

- The “mechanism”
- SoC-dependent
- The target frequency is decided elsewhere

struct cpufreq_driver

```
struct cpufreq_driver {  
    struct module *owner;  
    char name[CPUFREQ_NAME_LEN];  
  
    u8 flags;  
  
    int  (*init)    (struct cpufreq_policy *policy);  
    int  (*verify) (struct cpufreq_policy *policy);  
    ...  
};
```

struct cpufreq_driver

```
...
/* define one out of two */
int  (*setpolicy)(struct cpufreq_policy *policy);
int  (*target)   (struct cpufreq_policy *policy,
                  unsigned int target_freq,
                  unsigned int relation);

unsigned int (*get) (unsigned int cpu);
...
```

struct cpufreq_driver

```
...  
int    (*exit)    (struct cpufreq_policy *policy);  
int    (*suspend)(struct cpufreq_policy *policy);  
int    (*resume) (struct cpufreq_policy *policy);  
struct freq_attr  **attr;  
};
```

struct cpufreq_driver.flags

CPUFREQ_STICKY

- Driver remains even if `->init()` fails (typical)

CPUFREQ_CONST_LOOPS

- `loops_per_jiffy` doesn't change with `cpufreq`

.../mach-msm/cpufreq.c

```
struct cpufreq_driver msm_cpufreq_driver = {  
    .flags          = CPUFREQ_STICKY  
                    | CPUFREQ_CONST_LOOPS,  
    .init           = msm_cpufreq_init,  
    .verify         = msm_cpufreq_verify,  
    .target         = msm_cpufreq_target,  
    .get            = msm_cpufreq_get_freq,  
    .name           = "msm",  
    .attr           = msm_freq_attr,  
};
```

msm_cpufreq_init()

```
int msm_cpufreq_init(...)
{
    ...
    table = cpufreq_frequency_get_table(policy->cpu);
    ...
    /*
     * In 8625 both cpu core's frequency can not
     * be changed independently. Each cpu is bound to
     * same frequency. Hence set the cpumask to all cpu.
     */
    if (cpu_is_msm8625())
        cpumask_setall(policy->cpus);
    ...
}
```

msm_cpufreq_init()

```
...
if (cpufreq_frequency_table_cpuinfo(policy, table)) {
    policy->cpuinfo.min_freq = CONFIG_MSM_CPU_FREQ_MIN;
    policy->cpuinfo.max_freq = CONFIG_MSM_CPU_FREQ_MAX;
}
policy->min = CONFIG_MSM_CPU_FREQ_MIN;
policy->max = CONFIG_MSM_CPU_FREQ_MAX;
...
```

msm_cpufreq_init()

```
...
cur_freq = acpuclk_get_rate(policy->cpu);
if (cpufreq_frequency_table_target(policy,
    table, cur_freq,
    CPUFREQ_RELATION_H, &index) &&
    cpufreq_frequency_table_target(policy,
    table, cur_freq,
    CPUFREQ_RELATION_L, &index)) {
    return -EINVAL;
}
```

msm_cpufreq_init()

```
...
policy->cur = cur_freq;
policy->cpuinfo.transition_latency =
    acpuclk_get_switch_time() * NSEC_PER_USEC;

#ifdef CONFIG_SMP
    cpu_work = &per_cpu(cpufreq_work, policy->cpu);
    INIT_WORK(&cpu_work->work, set_cpu_work);
    init_completion(&cpu_work->complete);
#endif
...
```

msm_cpufreq_verify()

```
int msm_cpufreq_verify(struct cpufreq_policy *policy)
{
    cpufreq_verify_within_limits(policy,
        policy->cpuinfo.min_freq,
        policy->cpuinfo.max_freq);
    return 0;
}
```

msm_cpufreq_target()

```
int msm_cpufreq_target(struct cpufreq_policy *policy,
                       unsigned int target_freq,
                       unsigned int relation)
{
    ...
    table = cpufreq_frequency_get_table(policy->cpu);
    if (cpufreq_frequency_table_target(policy,
        table, target_freq, relation, &index)) {
        ret = -EINVAL;
        goto done;
    }
    ...
}
```

msm_cpufreq_target()

```
#ifdef CONFIG_SMP
    cpu_work = &per_cpu(cpufreq_work, policy->cpu);
    cpu_work->policy = policy;
    cpu_work->frequency = table[index].frequency;
    cpu_work->status = -ENODEV;

    ...
    ret = set_cpu_freq(cpu_work->policy,
                       cpu_work->frequency);
    ...
```


`struct cpufreq_policy`

Captures CPU frequency policy:

- What the target frequency is
- What the limits are
- The target frequency is decided elsewhere

struct cpufreq_policy

```
struct cpufreq_policy {  
    /* CPUs sharing clock, require sw coordination */  
    cpumask_var_t cpus;    /* Online CPUs only */  
    cpumask_var_t related_cpus; /* Online + Offline CPUs */  
    ...  
    unsigned int  cpu; /* CPU managing this policy */  
    ...  
};
```

struct cpufreq_policy

```
...
unsigned int  min;      /* in kHz */
unsigned int  max;      /* in kHz */
unsigned int  cur;      /* in kHz, for governors */
...
struct cpufreq_governor *governor;
};

DEFINE_PER_CPU(struct cpufreq_policy *, cpufreq_cpu_data);
```

struct cpufreq_governor

Defines the target CPU frequency:

- ... within policy limits
- ... according to a governing algorithm

Doesn't know how the target is achieved

struct cpufreq_governor

```
struct cpufreq_governor {
    char name[CPUFREQ_NAME_LEN];
    ...
    int (*governor)(struct cpufreq_policy *policy,
                    unsigned int event);
    ...
    struct list_head      governor_list;
    struct module         *owner;
};
```

struct cpufreq_governor

```
struct cpufreq_governor my_gov = {  
    .name      = "my-cpufreq-gov",  
    .governor  = my_cpufreq_gov,  
    .owner     = THIS_MODULE,  
};  
...  
cpufreq_register_governor(&my_gov);  
...
```

cpufreq_gov_performance

```
int cpufreq_gov_performance(...)
{
    switch (event) {
        case CPUFREQ_GOV_START:
        case CPUFREQ_GOV_LIMITS:
            __cpufreq_driver_target(policy, policy->max,
                                   CPUFREQ_RELATION_H);

            break;
        default:
            break;
    }
    return 0;
}
```

cpufreq_gov_performance

```
int __cpufreq_driver_target(...)
{
    ...
    if (target_freq > policy->max)
        target_freq = policy->max;
    ...
    if (target_freq == policy->cur)
        return 0;
    ...
    ret = cpufreq_driver->target(policy,
                                target_freq, relation);
    ...
}
```


cpufreq_gov_performance

```
int cpufreq_driver_target(...)
{
    policy = cpufreq_cpu_get(policy->cpu);
    ...
    ret = __cpufreq_driver_target(policy,
                                   target_freq, relation);
    ...
    cpufreq_cpu_put(policy);
    ...
}
```

cpufreq_gov_ondemand

```
void dbs_check_cpu(...)
{
    /* Extrapolated load of this CPU */
    unsigned int load_at_max_freq = 0;
    unsigned int avg_load_at_max_freq = 0;
    unsigned int max_load_freq;
    /* Current load across this CPU */
    unsigned int cur_load = 0;
    ...
}
```

cpufreq_gov_ondemand

```
...  
idle_time = get_cpu_idle_time(...);  
iowait_time = get_cpu_iowait_time(...);  
...
```