

Custom CPU Governor Example

Incorporating Battery Charge State

Bill Gatliff

`bgat@billgatliff.com`

Freelance Embedded Systems Developer

Functional Overview

Battery charge awareness:

- Decrease cpufreq when battery is low
- Increase available cpufreq when battery is charged
- Maximize cpufreq when connected to charger

Functional Overview

Caveats:

- No uniform battery state API
- What if we have more than one battery?
- What if we want to dial down features too?
- Is CPU really the largest power consumer?

That's why it's an “example” :-)

pm8921_bms_get_percent_charge()

```
struct pm8921_bms_chip *the_chip;

int pm8921_bms_get_percent_charge(void)
{
    if (!the_chip)
        return -EINVAL; /* uninitialized */
    return report_state_of_charge(the_chip);
}
EXPORT_SYMBOL_GPL(pm8921_bms_get_percent_charge);
```

batt_gov

```
struct cpufreq_governor batt_gov = {  
    .name      = "battery-pct",  
    .governor  = cpufreq_gov_batt,  
    .owner     = THIS_MODULE,  
};  
...  
cpufreq_register_governor(&batt_gov);  
...
```

batt_gov_info

```
struct batt_gov_info batt_gov_info = {  
    ...  
    struct cpufreq_policy *policy;  
    struct delayed_work work;  
    ...  
};  
struct batt_gov_info batt_gov_info;
```

cpufreq_gov_batt()

```
int cpufreq_gov_batt(...)
{
    switch (event) {
        case CPUFREQ_GOV_START:
        case CPUFREQ_GOV_LIMITS:
            batt_gov_info.policy = policy;
            INIT_DELAYED_WORK_DEFERRABLE(
                &batt_gov_info.work, do_batt_gov);
            schedule_delayed_work(&batt_gov_info.work, delay);
            break;
        ...
    }
    return 0;
}
```

do_batt_gov()

```
void do_batt_gov(struct work_struct *w)
{
    struct batt_gov_info *i = container_of(w, ...);
    struct cpufreq_policy *p = i->policy;
    int pct, pmax;

    pct = pm8921_bms_get_percent_charge();
    ...
}
```


do_batt_gov()

```
...
/* set policy max to a percentage of hardware max */
pmax = (p->cpuinfo.max - p->cpuinfo.min) * pct;
pmax = (pmax / 100) + p->cpuinfo.min;

__cpufreq_driver_target(i->policy, pmax,
                        CPUFREQ_RELATION_H);

schedule_delayed_work(w, delay);
}
```

do_batt_gov()

Still needs work:

- Synchronization!
- Missing or multiple batteries
- Is battery percentage the ONLY consideration?
- How can the user refine the policy settings?
- What about fully-discharged battery?

do_batt_gov()

```
void do_batt_gov(struct work_struct *w)
{
    ...
    if (pm8921_is_dc_chg_plugged_in()) {
        __cpufreq_driver_target(i->policy,
                                p->cpuinfo.max, CPUFREQ_RELATION_H);
        ...
    }
    ...
    schedule_delayed_work(w, delay);
}
```

do_batt_gov()

```
void do_batt_gov(struct work_struct *w)
{
    ...
    if (pmax < 30)
        pmax = 30;
    ...
}
```

do_batt_gov()

```
void do_batt_gov(struct work_struct *w)
{
    mutex_lock_interruptible(&i->mutex);
    ...
    __cpufreq_driver_target(i->policy, ...);
    ...
    schedule_delayed_work(w, delay);
    mutex_unlock(&i->mutex);
}
```

do_batt_gov()

```
int cpufreq_gov_batt(...)
{
    switch(event) {
        case CPUFREQ_GOV_STOP:
            cancel_delayed_work_sync(...);
            break;
    }
    ...
}
```

do_batt_gov()

```
void do_batt_gov(struct work_struct *w)
{
    ...
    if (!is_battery_present()) {
        pmax = p->cpuinfo.max; /* ? */
    }
    ...
}
```