

Power Management Quality-of-Service

Linux Power Management

Bill Gatliff

`bgat@billgatliff.com`

Freelance Embedded Systems Developer

What is “power management QOS”?

Quality-of-service:

- Defines minimum level of service required
- CPU interrupt latency or throughput
- Network latency
- ...

What is “power management QOS”?

“What’s the ‘power management’ part, then?”

- Guilt-by-association, mostly :-)

```
#include <linux/pm_qos.h>
```

Advisory Framework

It's mostly notifiers:

- Only platform subsystems can truly implement
- Very generic, by necessity

Three basic parameter classes:

- Latency
- Timeout (*)
- Throughput

Advisory Framework

```
struct notifier_block nb;
int my_callback(struct notifier_block *nb,
                unsigned long event,
                void *ignored)
{
    ...
}
```

Advisory Framework

```
/* add our notifier callback to the chain */  
nb.notifier_call = my_callback;  
notifier_chain_register(&chain, &nb);  
  
/* elsewhere, invoke the chain of callbacks */  
/* (i.e. when the event associated with the chain occurs) */  
notifier_call_chain(&chain, event, data);
```

Advisory Framework

Existing service classes:

- `PM_QOS_CPU_DMA_LATENCY`
- `PM_QOS_NETWORK_LATENCY`
- `PM_QOS_NETWORK_THROUGHPUT`

(It's fairly easy to create more as needed)

struct pm_qos_request

```
struct pm_qos_request {  
    struct plist_node node;  
    int pm_qos_class;  
    struct delayed_work work;  
};
```


pm_qos_add_request ()

Expresses an initial quality level:

- Use only once, e.g. during `probe ()`
- Used internally to “hook in” the associated notification

pm_qos_add_request ()

```
struct foo {  
    ...  
    struct pm_qos_request qos;  
    ...  
};
```

pm_qos_add_request ()

```
int foo_probe(...)
{
    ...
    /* usecs */
    latency = PM_QOS_CPU_DMA_LAT_DEFAULT_VALUE;
    pm_qos_add_request(&foo->qos,
        PM_QOS_CPU_DMA_LATENCY, latency);
    ...
}
```

pm_qos_update_request ()

Expresses a change in quality requirement:

- E.g. your device is at a different activity level
- Triggers a notification callback, as with the initial request

pm_qos_update_request()

```
int foo_open(...)
{
    ...
    pm_qos_update_request(&foo->qos, 100 /* usecs */);
    ...
}

int foo_close(...)
{
    ...
    pm_qos_update_request(&foo->qos, default_latency);
    ...
}
```

pm_qos_update_target ()

Callback to install a new quality target:

- For constraint implementers, not users
- Not normally invoked directly

pm_qos_update_target()

```
int pm_qos_update_target(...)  
{  
    ...  
    blocking_notifier_call_chain(...);  
    ...  
}
```

pm_qos_add_notifier()

Installes a notifier:

- Bridges gap between requestors, implementers
- Use this when your device can *assert* quality-of-service

drivers/cpuidle/cpuidle.c

```
int cpuidle_latency_notify(struct notifier_block *b,
                           unsigned long l, void *v)
{
    smp_call_function(smp_callback, NULL, 1);
    return NOTIFY_OK;
}

struct notifier_block cpuidle_latency_notifier = {
    .notifier_call = cpuidle_latency_notify,
};

void latency_notifier_init(struct notifier_block *n)
{
    pm_qos_add_notifier(PM_QOS_CPU_DMA_LATENCY, n);
}
```