

Completions

Basic Concurrency Features of the Linux Kernel

Bill Gatliff

`bgat@billgatliff.com`

Freelance Embedded Systems Developer

Completions

One-way communication between contexts:

- “The requested job is done”
- “An important event has occurred”

```
#include <linux/completion.h>
```

Completing with a semaphore (bad!)

```
foreground_task()  
{  
    sema_init(&sem);  
    start_external_task(&sem);  
    down(&sem);  
}  
  
interrupt_handler()  
{  
    up(sem);  
    return;  
}
```

Completions

Why?

- With semaphores, only one value blocks
- With completions, a block is anticipated

Semaphores are optimized for the opposite use case!

Completions

```
init_completion(struct completion *c)
```

- **Initializes** a struct completion

```
#define DECLARE_COMPLETION(name)
```

- **Declares and initializes**

Completions

```
void wait_for_completion(struct completion *c)
```

- Uninterruptible wait for completion

```
int wait_for_completion_interruptible(  
    struct completion *c)
```

- Interruptible wait for completion
- Returns nonzero if interrupted

Completions

```
unsigned long  
wait_for_completion_interruptible_timeout(  
    struct completion *c, unsigned long jiffies)
```

- Interruptible wait for completion, with timeout
- Returns `-ERESTARTSYS` on interruption
- Returns 0 on timeout
- Returns remaining jiffies on completion

Completions

```
void complete(struct completion *c)
```

- Signals completion
- Will wake up only one waiting context

```
void complete_all(struct completion *c)
```

- Signals completion for all waiting contexts

```
#define INIT_COMPLETION(struct completion c)
```

- Reinitializes after `complete_all()`

simple_completion.c

```
1  struct simple_completion_data {
2      struct file_operations fops;
3      struct cdev cdev;
4      struct completion compl;
5  };
6
7  static char stuff[32];
8  static int in, out;
9  #define NSTUFF (sizeof(stuff) / sizeof(*stuff))
```

simple_completion.c

```
1  static ssize_t
2  simple_completion_write (struct file *file, const char __user *buf,
3                          size_t count, loff_t *offp)
4  {
5      int new_in, orig_count = count;
6
7      while(count--> 0) {
8          new_in = in + 1;
9          if (new_in >= NSTUFF) new_in = 0;
10
11         if (new_in != out) {
12             copy_from_user(&stuff[new_in], buf++, 1);
13             in = new_in;
14         }
15         else break;
16     }
17
18     complete(&simple_completion_data.compl);
19     return orig_count;
20 }
```

simple_completion.c

```
1  static ssize_t
2  simple_completion_read (struct file *file, char __user *buf,
3                          size_t count, loff_t *offp)
4  {
5      int bytes = 0;
6
7      if (wait_for_completion_interruptible(&simple_completion_data.compl))
8          return -EINTR;
9
10     while (out != in) {
11         copy_to_user(buf++, &stuff[out], 1);
12         bytes++;
13         if (++out >= NSTUFF) out = 0;
14     }
15
16     return bytes;
17 }
```

simple_completion_wrong.c

Question:

- What if we move the `complete()` call?

simple_completion_wrong.c

```
1  static ssize_t
2  simple_completion_write (struct file *file, const char __user *buf,
3                          size_t count, loff_t *offp)
4  {
5      int new_in, orig_count = count;
6
7      while(count--) {
8          new_in = in + 1;
9          if (new_in >= NSTUFF) new_in = 0;
10
11         if (new_in != out) {
12             copy_from_user(&stuff[new_in], buf++, 1);
13             complete(&simple_completion_data.compl);
14             in = new_in;
15         }
16         else break;
17     }
18
19     return orig_count;
20 }
```

kthread-completion.c

```
1  static int kthrd(void *arg)
2  {
3      struct sensor *s = arg;
4      s32 ret;
5
6      while (!kthread_should_stop()) {
7          init_completion(&s->complete);
8          wait_for_completion_interruptible(&s->complete);
9
10         ret = i2c_smbus_read_byte(&s->client);
11         if (ret >= 0)
12             pr_err("%s: i2c_smbus_read_byte returned %x\n", __func__, ret);
13     }
14     return 0;
15 }
16
17 static irqreturn_t sensor_interrupt(int irq, void *data)
18 {
19     struct sensor *s = data;
20     complete(&s->complete);
21     return IRQ_HANDLED;
22 }
```

Completions

Basic Concurrency Features of the Linux Kernel

Bill Gatliff

`bgat@billgatliff.com`

Freelance Embedded Systems Developer